

CSC139 Operating System Principles

Fall 2019, Part 4-3

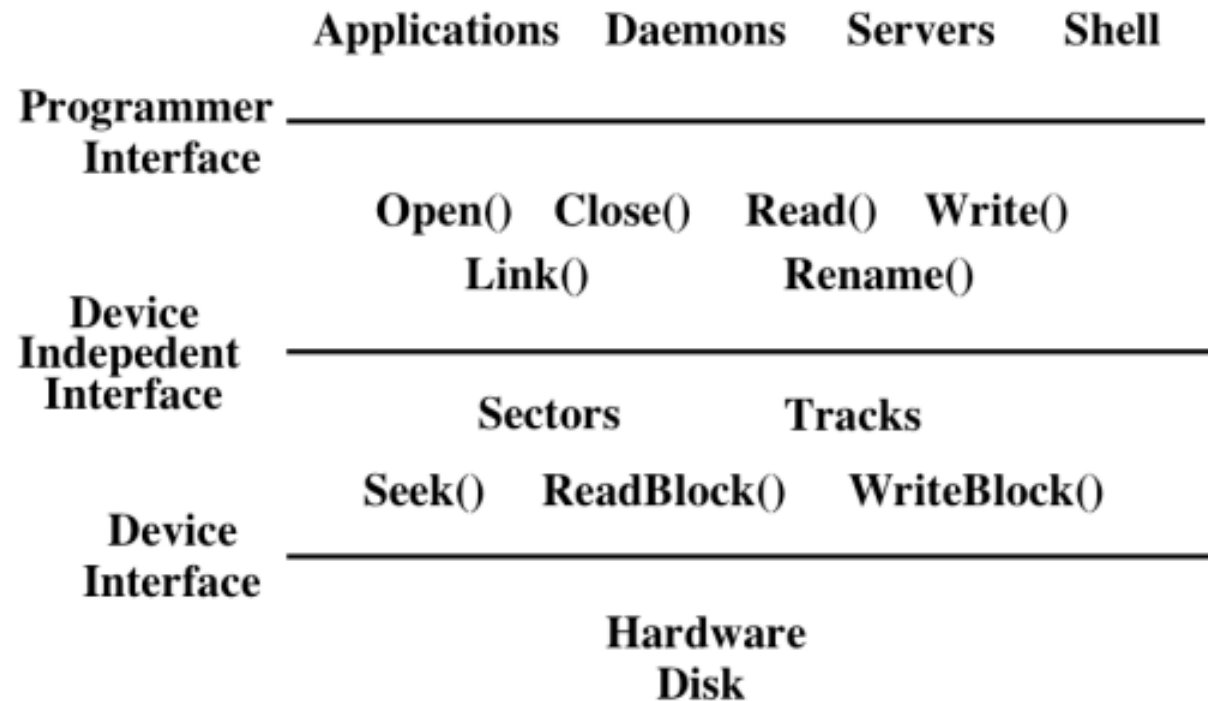
Instructor: Dr. Yuan Cheng

Today: File System Functionality

- Remember the high-level view of the OS as a translator from the user abstraction to the hardware reality

User Abstraction		Hardware Reality
Processes/Threads	<=OS=>	CPU
Address Space	<=OS=>	Memory
Files	<=OS=>	Disk

File System Abstraction



User Requirements on Data

- **Persistence:** data stays around between jobs, power cycles, crashes
- **Speed:** can get to data quickly
- **Size:** can store lots of data
- **Sharing/Protection:** users can share data where appropriate or keep it private when appropriate
- **Ease of Use:** user can easily find, examine, modify, etc. data

Hardware/OS Features

- Hardware provides:
 - Persistence: Disks provide non-volatile memory
 - Speed: Speed gained through random access
 - Size: Disks keep getting bigger (typical disk on a PC = 500GB)
- OS provides:
 - Persistence: redundancy allows recovery from some additional failures
 - Sharing/Protection: Unix provides read, write, execute privileges for files
 - Ease of Use
 - Associating names with chunks of data (files)
 - Organize large collections of files into directories
 - Transparent mapping of the user's concept of files and directories onto locations on disks
 - Search facility in file systems (Spotlight in Mac OS X)

Files

- File: Logical unit of data on a storage device
 - Formally, named collection of related information recorded on secondary storage
 - Example: reader.cc, a.out
- Files can contain programs (source, binary) or data
- Files can be structured or unstructured
 - Unix implements files as a series of bytes (unstructured)
 - IBM mainframes implements files as a series of records or objects (structured)
- File attributes: name, type, location, size, protection, creation time

File Names

- Three types of names (abstractions)
 - inode (low-level names)
 - path (human readable)
 - file descriptor (runtime state)

Inodes

- Each file has exactly one inode number
- Inodes are unique (at a given time) within a FS
- Numbers may be recycled after deletes
- Show inodes via `stat`
 - `$ stat <file or dir>`

Path (multiple directories)

- A directory is a file
 - Associated with an inode
- Contains a list of <user-readable name, low-level name> pairs

User Interface to the File System

- Common file operations:
 - Data operations:
 - Create(), Delete(), Open(), Close(), Read(), Write(), Seek()
 - Naming operations:
 - HardLink(), SoftLink(), Rename(), SetAttribute(), GetAttribute()

OS File Data Structures

1. Open file table - shared by all processes with an open file.
 - open count
 - file attributes, including ownership, protection information, access times, ...
 - location(s) of file on disk
 - pointers to location(s) of file in memory
2. Per-process file table
 - for each file,
 - pointer to entry in the open file table
 - current position in file (offset)
 - mode in which the process will access the file (r, w, rw)
 - pointers to file buffer

File Operations: Creating a File

- Create(name)
 - Allocate disk space (check disk quotas, permissions, etc.)
 - Create a file descriptor for the file including name, location on disk, and all file attributes.
 - Add the file descriptor to the directory that contains the file.
 - Optional file attribute: file type (Word file, executable, etc.)

File Operations: Deleting a File

- Delete(name)
 - Find the directory containing the file.
 - Free the disk blocks used by the file.
 - Remove the file descriptor from the directory.
 - Behavior dependent on hard links

File Operations: Open and Close

- `fileId = Open(name, mode)`
 - Check if the file is already open by another process. If not,
 - Find the file.
 - Copy the file descriptor into the system-wide open file table.
 - Check the protection of the file against the requested mode. If not ok, abort
 - Increment the open count.
 - Create an entry in the process's file table pointing to the entry in the systemwide file table. Initialize the current file pointer to the start of the file.
- `Close(fileId)`
 - Remove the entry for the file in the process's file table.
 - Decrement the open count in the system-wide file table.
 - If the open count == 0, remove the entry in the system-wide file table.

File Operations: Reading a File

- Read(fileID, from, size, bufAddress)
 - random/direct access – OS reads “size” bytes from file position “from” into “bufAddress”

```
for (i = from; i < from + size; i++)  
    bufAddress[i - from] = file[i];
```
- Read(fileID, size, bufAddress) - sequential access
 - OS reads “size” bytes from current file position, fp, into “bufAddress” and increments current file position by size

```
for (i = 0; i < size; i++)  
    bufAddress[i] = file[fp + i];  
fp += size;
```

File Operations

- **Write** is similar to reads, but copies from the buffer to the file.
- **Seek** just updates fp.
- **Memory mapping** a file
 - Map a part of the portion virtual address space to a file
 - Read/write to that portion of memory \implies OS reads/writes from corresponding location in the file
 - File accesses are greatly simplified (no read/write call are necessary)

File Access Methods


- Common file access patterns from the programmer's perspective
 - Sequential: data processed in order, a byte or record at a time.
 - Most programs use this method
 - Example: compiler reading a source file.
 - Direct: address a block based on a key value.
 - Example: database search, hash table, dictionary
- Common file access patterns from the OS perspective:
 - Sequential: keep a pointer to the next byte in the file. Update the pointer on each read/write.
 - Direct: address any block in the file directly given its offset within the file.

Naming and Directories

- Need a method of getting back to files that are left on disk.
- OS uses numbers for each files
 - Users prefer textual names to refer to files.
 - Directory: OS data structure to map names to file descriptors
- Naming strategies
 - Single-Level Directory: One name space for the entire disk, every name is unique.
 1. Use a special area of disk to hold the directory.
 2. Directory contains pairs.
 3. If one user uses a name, no one else can.
 4. Some early computers used this strategy. Early personal computers also used this strategy because their disks were very small.
 - Two Level Directory: each user has a separate directory, but all of each user's files must still have unique names

Single-Level Directory

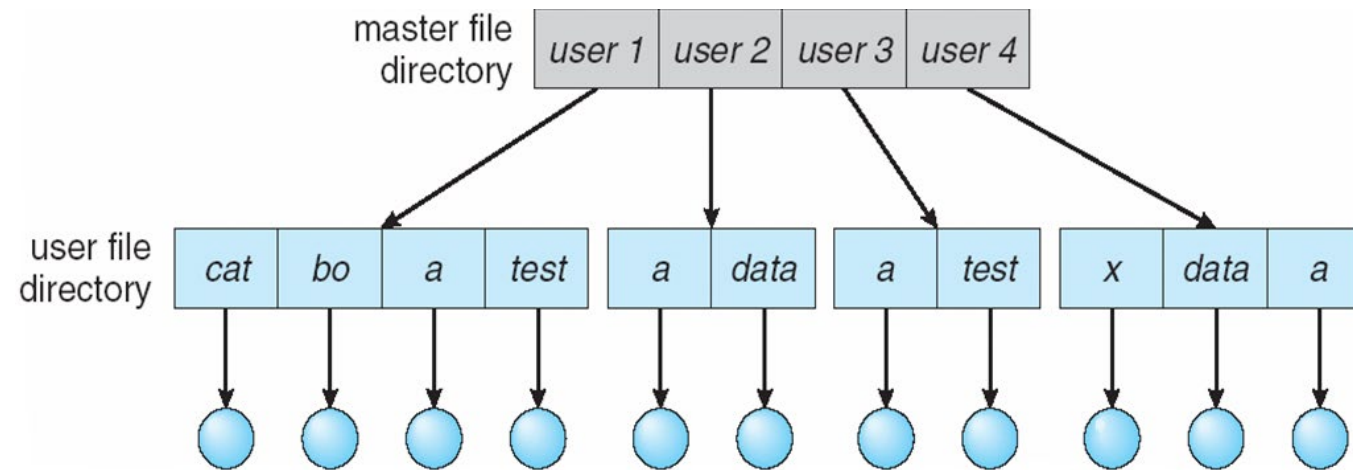
- A single directory for all users



- Naming problem
- Grouping problem

Two-Level Directory

- Separate directory for each user

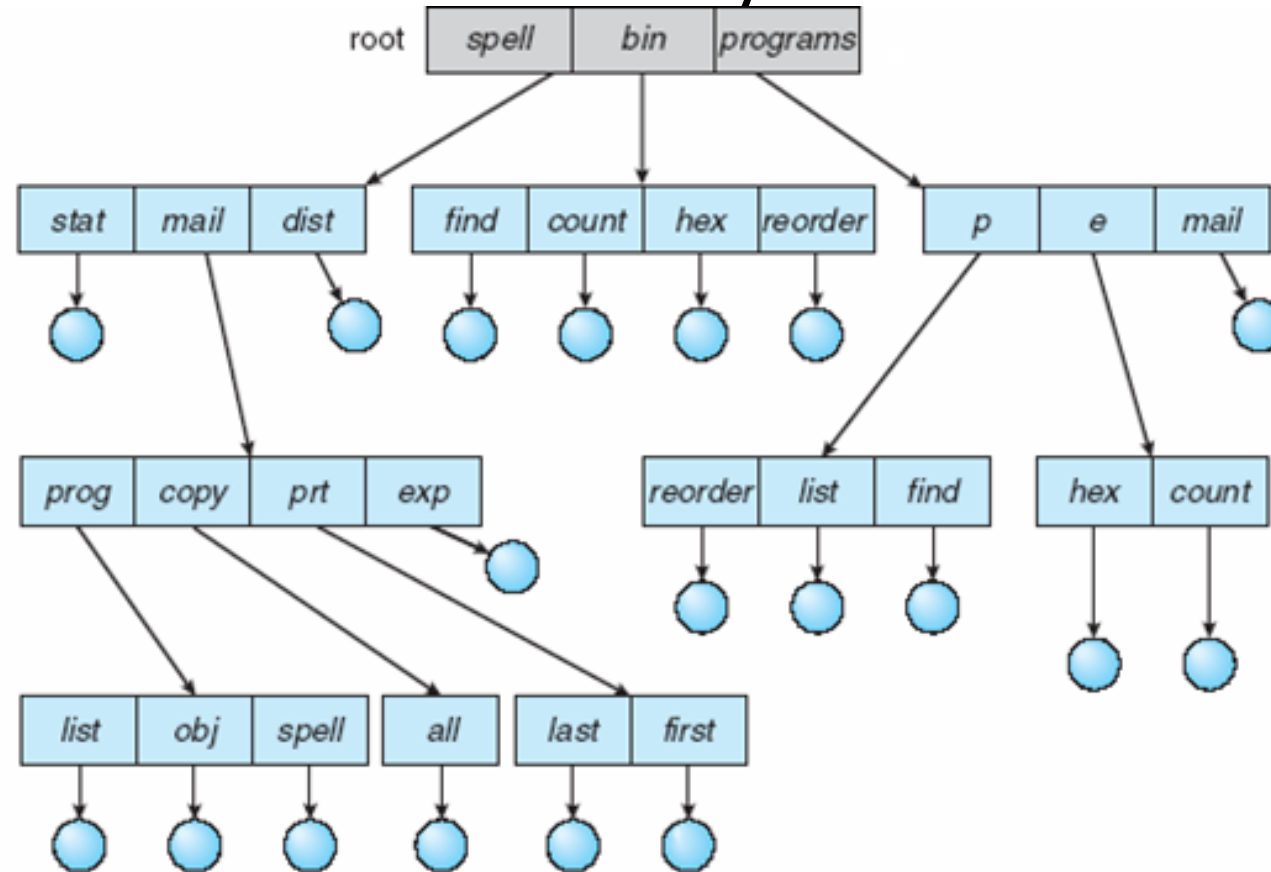


- Path name
- Can have the same file name for different user
- Efficient searching
- No grouping capability

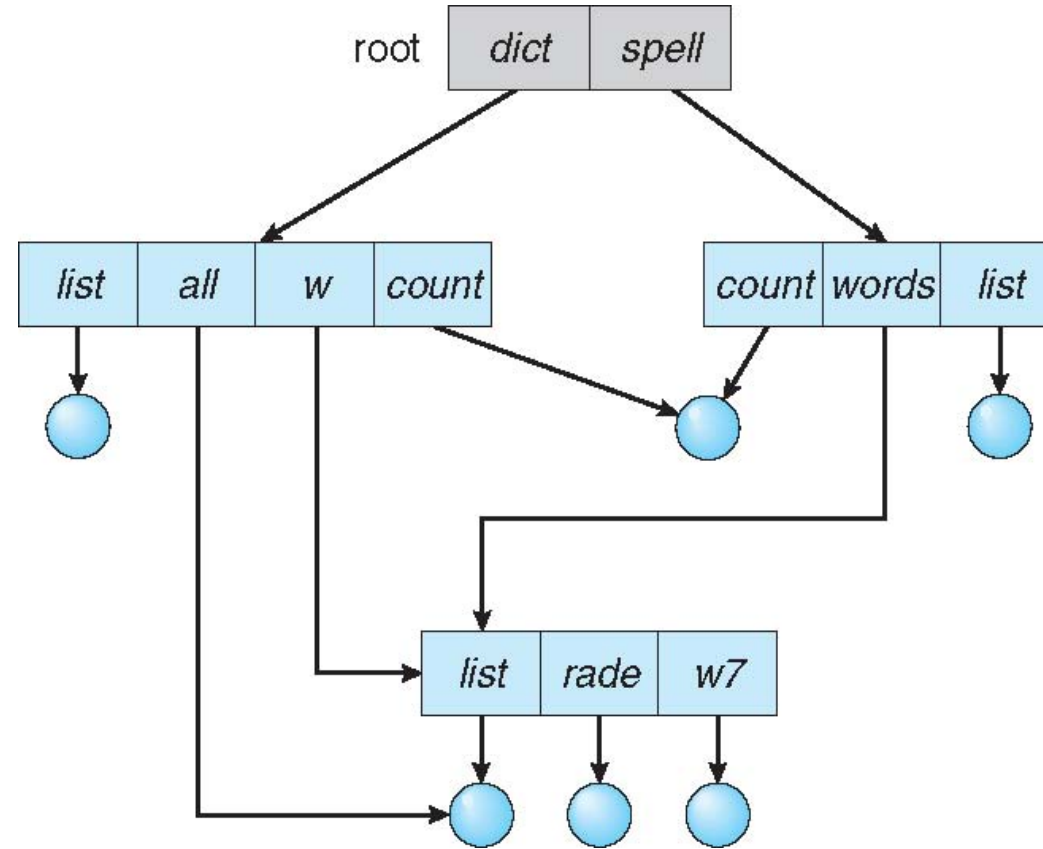
Naming Strategies (cont.)

- Multilevel Directories - tree structured name space (Unix, and all other modern operating systems).
 1. Store directories on disk, just like files except the file descriptor for directories has a special flag bit.
 2. User programs read directories just like any other file, but only special system calls can write directories.
 3. Each directory contains pairs in no particular order. The file referred to by a name may be another directory.
 4. There is one special root directory. Example: How do we look up name: /usr/bin/l
- Limitations with basic tree structure
 - Difficult to share file across directories and users
 - Can't have multiple file names

Tree-Structured Directory



Acyclic-Graph Directory



Referential Naming

- Hard links (Unix: *ln* command)
 - A hard link adds a second connection to a file
 - Example: creating a hard link from B to A

Initially:	A → file #100
<hr/>	
After " <i>ln</i> A B":	A → file #100 B → file #100

- OS maintains reference counts, so it will only delete a file after the last link to it has been deleted.

Referential Naming (cont.)

- Soft links (Unix: `ln -s` command)
 - A soft link only makes a symbolic pointer from one file to another.
 - Example: creating a soft link from B to A

Initially:	$A \rightarrow \text{file \#100}$
<hr/>	
After " <code>ln A B</code> ":	$A \rightarrow \text{file \#100}$ $B \rightarrow A$

- removing B does not affect A
- removing A leaves the name B in the directory, but its contents no longer exists

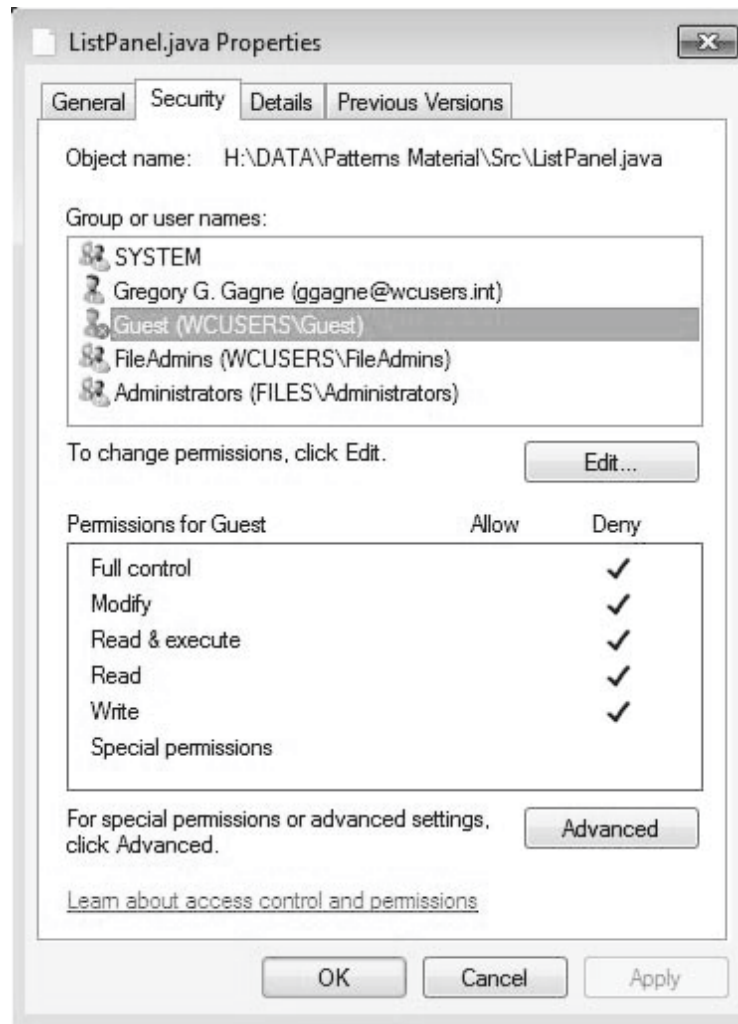
Directory Operations

- Search for a file: locate an entry for a file
- Create a file: add a directory listing
- Delete a file: remove directory listing
- List a directory: list all files (ls command in UNIX)
- Rename a file
- Traverse the file system

Protection

- The OS must allow users to control sharing of their files => control access to files
- Grant or deny access to file operations depending on protection information
- Access lists and groups (Windows NT)
 - Keep an access list for each file with user name and type of access
 - Lists can become large and tedious to maintain
- Access control bits (UNIX)
 - Three categories of users (owner, group, world)
 - Three types of access privileges (read, write, execute)
 - Maintain a bit for each combination (111101000 = rwxr-x---)

Windows 7 Access-Control List Management



A Sample UNIX Directory Listing

-rw-rw-r--	1 pbg	staff	31200	Sep 3 08:30	intro.ps
drwx-----	5 pbg	staff	512	Jul 8 09:33	private/
drwxrwxr-x	2 pbg	staff	512	Jul 8 09:35	doc/
drwxrwx---	2 pbg	student	512	Aug 3 14:13	student-proj/
-rw-r--r--	1 pbg	staff	9423	Feb 24 2003	program.c
-rwxr-xr-x	1 pbg	staff	20471	Feb 24 2003	program
drwx--x--x	4 pbg	faculty	512	Jul 31 10:31	lib/
drwx-----	3 pbg	staff	1024	Aug 29 06:52	mail/
drwxrwxrwx	3 pbg	staff	512	Jul 8 09:35	test/

Exit Slips

- Take 1-2 minutes to reflect on this lecture
- On a sheet of paper write:
 - One thing you learned in this lecture
 - One thing you didn't understand

Next class

- We will continue to discuss:
 - File Systems
- Reading assignment:
 - SGG: Ch. 14

Acknowledgment

- The slides are partially based on the ones from
 - The book site of *Operating System Concepts (Tenth Edition)*: <http://os-book.com/>