

# CSC 177: Data Warehousing and Data Mining

## Project 1: Data Preprocessing

Jagan Chidella

Group: An Lam, Jimmy Le, Tom Amir,  
Dianna Melendez, Amrit Singh, Talal Jawaaid, Min Li

## Reader Manual:

Starting from Page 3 to Page 25 are screenshots of the code given in Tutorial 4 for Data Preprocessing. Each method of Preprocessing given to us was analyzed for our personal understanding. After realizing the capabilities of the various methods of Data Preprocessing, we began to seek out our data sets online.

After searching online, we chose two data sets that were interesting that we were going to thoroughly analyze for Part 3B of our project:

- 1) [Drugs](#): In this zip file, we are using the “Product” Excel Sheet. We made sure to save it as a “.csv” file before proceeding applying the Pre-Processing techniques that we learned.
- 2) [London Air](#): In this data set, we found that we were able to apply the more interesting Pre-Processing techniques, such as: missing values, and saving a dataframe.

Since the purpose of this project is Pre-Processing, we wanted to take data sets and clean them up in such a way where it was an easily readable format. In addition to being easily read, we wanted less attributes to deal with, while dropping unnecessary data. Starting from Page 17, until Page 23, the techniques that we applied were: Missing Values, Outliers, Duplicate Data, Shuffling Dataframes, Sorting Dataframes, Saving a Dataframes, Dropping Fields, Calculated Fields, Feature Normalization, Concatenation, Aggregation, and PCA.

After Pre-Processing our data, we ran a thorough analysis covered in Pages 24 onward. We split the data into training and testing by the ratio of 80:20. We figured out the mean and the standard deviation of each respective set and drew a conclusion from it.

Starting here, we ran through every method of Pre-Processing covered in the tutorial and posted the screenshots of the outputs from each, and then analyzed the results.

**1. Data Quality Issues:** The quality of the data is a major concern so we start by displaying our data. We start by printing out the number of rows, and the number of columns in the data set.

```
In [1]: import pandas as pd
data = pd.read_csv('https://archive.ics.uci.edu/ml/machine-learning-databases/breast-cancer-wisconsin/breast-cancer-wisconsin.data')
data.columns = ['Sample code', 'Clump Thickness', 'Uniformity of Cell Size', 'Uniformity of Cell Shape', 'Marginal Adhesion', 'Single Epithelial Cell Size', 'Bare Nuclei', 'Bland Chromatin', 'Normal Nucleoli', 'Mitoses', 'Class']

data = data.drop(['Sample code'],axis=1)
print('Number of instances = %d' % (data.shape[0]))
print('Number of attributes = %d' % (data.shape[1]))
data.head()
```

Number of instances = 699  
Number of attributes = 10

Out[1]:

	Clump Thickness	Uniformity of Cell Size	Uniformity of Cell Shape	Marginal Adhesion	Single Epithelial Cell Size	Bare Nuclei	Bland Chromatin	Normal Nucleoli	Mitoses	Class
0	5	1	1	1	2	1	3	1	1	2
1	5	4	4	5	7	10	3	2	1	2
2	3	1	1	1	2	2	3	1	1	2
3	6	8	8	1	3	4	3	7	1	2
4	4	1	1	3	2	1	3	1	1	2

**2a. Missing Values:** In the data set, it is important to know which columns have data that is missing in it. In this section of code, each column is parsed, counting each missing value.

```
In [2]: import numpy as np

data = data.replace('?',np.NaN)

print('Number of instances = %d' % (data.shape[0]))
print('Number of attributes = %d' % (data.shape[1]))

print('Number of missing values:')
for col in data.columns:
    print('\t%s: %d' % (col,data[col].isna().sum()))
```

Number of instances = 699  
Number of attributes = 10  
Number of missing values:  
Clump Thickness: 0  
Uniformity of Cell Size: 0  
Uniformity of Cell Shape: 0  
Marginal Adhesion: 0  
Single Epithelial Cell Size: 0  
Bare Nuclei: 16  
Bland Chromatin: 0  
Normal Nucleoli: 0  
Mitoses: 0  
Class: 0

**2b. Missing Values:** After pinpointing which column had missing values, the rows that had missing values were replaced with the median of the set in order to prevent skewed data.

```
In [3]: data2 = data['Bare Nuclei']

print('Before replacing missing values:')
print(data2[20:25])
data2 = data2.fillna(data2.median())

print('\nAfter replacing missing values:')
print(data2[20:25])

Before replacing missing values:
20    10
21     7
22     1
23    NaN
24     1
Name: Bare Nuclei, dtype: object

After replacing missing values:
20    10
21     7
22     1
23     1
24     1
Name: Bare Nuclei, dtype: object
```

**2c. Missing Values:** Another approach to missing values is to simply drop the whole row if a column has missing data in it.

```
In [4]: print('Number of rows in original data = %d' % (data.shape[0]))

data2 = data.dropna()
print('Number of rows after discarding missing values = %d' % (data2.shape[0]))

Number of rows in original data = 699
Number of rows after discarding missing values = 683
```

**3a. Outliers:** Original data provided

```
In [1]: import pandas as pd
data = pd.read_csv('https://archive.ics.uci.edu/ml/machine-learning-databases/breast-cancer-wisconsin/breast-cancer-wisconsin.data')
data.columns = ['Sample code', 'Clump Thickness', 'Uniformity of Cell Size', 'Uniformity of Cell Shape', 'Marginal Adhesion', 'Single Epithelial Cell Size', 'Bare Nuclei', 'Bland Chromatin', 'Normal Nucleoli', 'Mitoses', 'Class']

data = data.drop(['Sample code'],axis=1)
print('Number of instances = %d' % (data.shape[0]))
print('Number of attributes = %d' % (data.shape[1]))
data.head()
```

Number of instances = 699  
Number of attributes = 10

```
Out[1]:
```

	Clump Thickness	Uniformity of Cell Size	Uniformity of Cell Shape	Marginal Adhesion	Single Epithelial Cell Size	Bare Nuclei	Bland Chromatin	Normal Nucleoli	Mitoses	Class
0	5	1	1	1	2	1	3	1	1	2
1	5	4	4	5	7	10	3	2	1	2
2	3	1	1	1	2	2	3	1	1	2
3	6	8	8	1	3	4	3	7	1	2
4	4	1	1	3	2	1	3	1	1	2

### 3b. Outliers: Calculating the Z-Score to help locate outliers

```
In [6]: Z = (data2-data2.mean())/data2.std()  
Z[20:25]
```

Out[6]:

	Clump Thickness	Uniformity of Cell Size	Uniformity of Cell Shape	Marginal Adhesion	Single Epithelial Cell Size	Bare Nuclei	Bland Chromatin	Normal Nucleoli	Mitoses
20	0.917080	-0.044070	-0.406284	2.519152	0.805662	1.771569	0.640688	0.371049	1.405526
21	1.982519	0.611354	0.603167	0.067638	1.257272	0.948266	1.460910	2.335921	-0.343666
22	-0.503505	-0.699494	-0.742767	-0.632794	-0.549168	-0.698341	-0.589645	-0.611387	-0.343666
23	1.272227	0.283642	0.603167	-0.632794	-0.549168	NaN	1.460910	0.043570	-0.343666
24	-1.213798	-0.699494	-0.742767	-0.632794	-0.549168	-0.698341	-0.179534	-0.611387	-0.343666

### 3c. Outliers: Discarding the outliers. Significantly reduces the number of rows

```
In [7]: print('Number of rows before discarding outliers = %d' % (Z.shape[0]))  
  
Z2 = Z.loc[((Z > -3).sum(axis=1)==9) & ((Z <= 3).sum(axis=1)==9),:]  
print('Number of rows after discarding missing values = %d' % (Z2.shape[0]))  
  
Number of rows before discarding outliers = 699  
Number of rows after discarding missing values = 632
```

### 4a. Duplicate Data

Some datasets, especially when obtained by merging multiple data sources can contain duplicate sets. The code below will check for duplicate data.

```
In [8]: dups = data.duplicated()  
print('Number of duplicate rows = %d' % (dups.sum()))  
data.loc[[11,28]]
```

Number of duplicate rows = 236

Out[8]:

	Clump Thickness	Uniformity of Cell Size	Uniformity of Cell Shape	Marginal Adhesion	Single Epithelial Cell Size	Bare Nuclei	Bland Chromatin	Normal Nucleoli	Mitoses	Class
11	2	1	1	1	2	1	2	1	1	2
28	2	1	1	1	2	1	2	1	1	2

### 4b. Duplicate Data

Some datasets, especially when obtained by merging multiple data sources can contain duplicate sets. The code below will discard the duplicate data after duplicate data is found.

```
In [9]: print('Number of rows before discarding duplicates = %d' % (data.shape[0]))  
data2 = data.drop_duplicates()  
print('Number of rows after discarding duplicates = %d' % (data2.shape[0]))  
  
Number of rows before discarding duplicates = 699  
Number of rows after discarding duplicates = 463
```

## 5. Shuffling Dataframes

The code below will shuffle the data frame.

```
In [23]: import os
import numpy as np
import pandas as pd

path = "./downloads/"

filename_read = os.path.join(path, "auto-mpg.csv")
df = pd.read_csv(filename_read, na_values=['NA', '?'])

#np.random.seed(30) # Uncomment this line to get the same shuffle each time

df = df.reindex(np.random.permutation(df.index))
df.reset_index(inplace=True, drop=True)
# use inplace=False
df
```

```
Out[23]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	origin	name
0	24.0	4	140.0	92.0	2865	16.4	82	1	ford fairmont futura
1	18.0	6	250.0	88.0	3021	16.5	73	1	ford maverick
2	37.2	4	86.0	65.0	2019	16.4	80	3	mitsubishi 310
3	30.9	4	105.0	75.0	2230	14.5	78	1	dodge omni
4	20.0	6	198.0	95.0	3102	16.5	74	1	plymouth duster
5	33.8	4	97.0	67.0	2145	18.0	80	3	subaru dl

## 6. Sorting Dataframes

The code below sorts data frame by name and will show in ascending order.

```
In [24]: df = df.sort_values(by='name', ascending=True)
df
```

```
Out[24]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	origin	name
42	13.0	8	360.0	175.0	3821	11.0	73	1	amc ambassador brougham
263	15.0	8	390.0	190.0	3850	8.5	70	1	amc ambassador dpl
40	17.0	8	304.0	150.0	3672	11.5	72	1	amc ambassador sst
144	19.4	6	232.0	90.0	3210	17.2	78	1	amc concord
296	24.3	4	151.0	90.0	3003	20.1	80	1	amc concord
204	18.1	6	258.0	120.0	3410	15.1	78	1	amc concord d/i
48	23.0	4	151.0	NaN	3035	20.5	82	1	amc concord dl
147	20.2	6	232.0	90.0	3265	18.2	79	1	amc concord dl 6
384	18.0	6	232.0	100.0	2789	15.0	73	1	amc gremlin
213	19.0	6	232.0	100.0	2634	13.0	71	1	amc gremlin
69	20.0	6	232.0	100.0	2914	16.0	75	1	amc gremlin
22	21.0	6	199.0	90.0	2648	15.0	70	1	amc gremlin
281	18.0	6	232.0	100.0	2945	16.0	73	1	amc hornet
184	19.0	6	232.0	100.0	2901	16.0	74	1	amc hornet
354	18.0	6	199.0	97.0	2774	15.5	70	1	amc hornet
67	22.5	6	232.0	90.0	3085	17.6	76	1	amc hornet

```
In [25]: print("The first car is: {}".format(df['name'].iloc[0]))

The first car is: amc ambassador brougham
```

```
In [25]: print("The first car is: {}".format(df['name'].iloc[0]))

The first car is: amc ambassador brougham
```

## 7. Saving a Dataframes

The code below will shuffle and then save a new copy to the data.

```
In [27]: import os
import pandas as pd
import numpy as np

path = "./downloads/"

filename_read = os.path.join(path, "auto-mpg.csv")
filename_write = os.path.join(path, "auto-mpg-shuffle.csv")
df = pd.read_csv(filename_read, na_values=['NA', '?'])
df = df.reindex(np.random.permutation(df.index))
df.to_csv(filename_write, index=False) # Specify index = false to not write row numbers
print("Done")

Done
```

## 8. Dropping Fields

When fields are of no value to the data set we can drop them. The code below removes the column from the MPG dataset.

```
In [28]: import os
import pandas as pd
import numpy as np

path = "./downloads/"

filename_read = os.path.join(path, "auto-mpg.csv")
df = pd.read_csv(filename_read, na_values=['NA', '?'])

print("Before drop: {}".format(df.columns))
df.drop('name', axis=1, inplace=True)
print("After drop: {}".format(df.columns))
df[0:5]

Before drop: Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
'acceleration', 'year', 'origin', 'name'],
dtype='object')
After drop: Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
'acceleration', 'year', 'origin'],
dtype='object')
```

```
Out[28]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	origin
0	18.0	8	307.0	130.0	3504	12.0	70	1
1	15.0	8	350.0	165.0	3693	11.5	70	1
2	18.0	8	318.0	150.0	3436	11.0	70	1
3	16.0	8	304.0	150.0	3433	12.0	70	1
4	17.0	8	302.0	140.0	3449	10.5	70	1

## 9. Calculated Fields

Calculated fields can be added as new fields to the data frame that are calculated from our data. We can create a new column and give the weight in kilograms.

```
In [28]: import os
import pandas as pd
import numpy as np

path = "./downloads/"

filename_read = os.path.join(path, "auto-mpg.csv")
df = pd.read_csv(filename_read, na_values=['NA', '?'])

print("Before drop: {}".format(df.columns))
df.drop('name', axis=1, inplace=True)
print("After drop: {}".format(df.columns))
df[0:5]
```

Before drop: Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight', 'acceleration', 'year', 'origin', 'name'], dtype='object')

After drop: Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight', 'acceleration', 'year', 'origin'], dtype='object')

```
Out[28]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	origin
0	18.0	8	307.0	130.0	3504	12.0	70	1
1	15.0	8	350.0	165.0	3693	11.5	70	1
2	18.0	8	318.0	150.0	3436	11.0	70	1
3	16.0	8	304.0	150.0	3433	12.0	70	1
4	17.0	8	302.0	140.0	3449	10.5	70	1

## 10. Feature Normalization

Feature normalization data allows us to put numbers in standard form so that two values can be compared. This code replaces the mpg with a z-score. Cars with average MPG will be near zero is above average. And below zero is below average. Outliers at >3 or <-3

```
In [30]: import os
import pandas as pd
import numpy as np
from scipy.stats import zscore

path = "./downloads/"

filename_read = os.path.join(path, "auto-mpg.csv")
df = pd.read_csv(filename_read, na_values=['NA', '?'])
df['mpg'] = zscore(df['mpg'])
df
```

```
Out[30]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	origin	name
0	-0.706439	8	307.0	130.0	3504	12.0	70	1	chevrolet chevelle malibu
1	-1.090751	8	350.0	165.0	3693	11.5	70	1	buick skylark 320
2	-0.706439	8	318.0	150.0	3436	11.0	70	1	plymouth satellite
3	-0.962647	8	304.0	150.0	3433	12.0	70	1	amc rebel sst
4	-0.834543	8	302.0	140.0	3449	10.5	70	1	ford torino
5	-1.090751	8	429.0	198.0	4341	10.0	70	1	ford galaxie 500
6	-1.218855	8	454.0	220.0	4354	9.0	70	1	chevrolet impala
7	-1.218855	8	440.0	215.0	4312	8.5	70	1	plymouth fury iii
8	-1.218855	8	455.0	225.0	4425	10.0	70	1	pontiac catalina
9	-1.090751	8	390.0	190.0	3850	8.5	70	1	amc ambassador dpl



## 12. Concatenating Rows and Columns

This data shows how rows and columns can be concatenated together to form new data frames.

```
In [32]: #Create a new dataframe from name and horsepower

import os
import pandas as pd
import numpy as np
from scipy.stats import zscore

path = "./downloads/"

filename_read = os.path.join(path, "auto-mpg.csv")
df = pd.read_csv(filename_read, na_values=['NA', '?'])
col_horsepower = df['horsepower']
col_name = df['name']
result = pd.concat([col_name, col_horsepower], axis=1)
result
```

```
Out[32]:
```

	name	horsepower
0	chevrolet chevelle malibu	130.0
1	buick skylark 320	165.0
2	plymouth satellite	150.0
3	amc rebel sst	150.0
4	ford torino	140.0
5	ford galaxie 500	198.0
6	chevrolet impala	220.0
7	plymouth fury iii	215.0

## 13. Tensorflow

- Tensorflow allows you to build the feature vector in the format expected from raw data.
  1. Encoding Data
    - encode\_text\_dumy
    - Encode\_text\_index
  2. Normalizing Data
    - Encode\_numeric\_zscore
    - Encode\_numeric\_range
  3. Dealing with missing data:
    - missing\_median
  4. Removing outliers:
    - Remove\_outliers
  5. Create feature vector and target vector
    - To\_xy
  6. Other utility functions:
    - Hms\_string
    - chart\_regression

## 14. Label encoding, one hot encoding, creating X/Y for TensorFlow

```
In [5]: df=pd.read_csv("data/iris.csv",na_values=['NA','?'])  
df
```

Out[5]:

	sepal_l	sepal_w	petal_l	petal_w	species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa
5	5.4	3.9	1.7	0.4	Iris-setosa
6	4.6	3.4	1.4	0.3	Iris-setosa
7	5.0	3.4	1.5	0.2	Iris-setosa
8	4.4	2.9	1.4	0.2	Iris-setosa
9	4.9	3.1	1.5	0.1	Iris-setosa
10	5.4	3.7	1.5	0.2	Iris-setosa
11	4.8	3.4	1.6	0.2	Iris-setosa
12	4.8	3.0	1.4	0.1	Iris-setosa
13	4.3	3.0	1.1	0.1	Iris-setosa
14	5.8	4.0	1.2	0.2	Iris-setosa
15	5.7	4.4	1.5	0.4	Iris-setosa
16	5.4	3.9	1.3	0.4	Iris-setosa
17	5.1	3.5	1.4	0.3	Iris-setosa
18	5.7	3.8	1.7	0.3	Iris-setosa
19	5.1	3.8	1.5	0.3	Iris-setosa
20	5.4	3.4	1.7	0.2	Iris-setosa
21	5.1	3.7	1.5	0.4	Iris-setosa
22	4.6	3.6	1.0	0.2	Iris-setosa
23	5.1	3.3	1.7	0.5	Iris-setosa
24	4.8	3.4	1.9	0.2	Iris-setosa
25	5.0	3.0	1.6	0.2	Iris-setosa

## 15. Encoding labels before calling to\_xy()

```
In [8]: df=pd.read_csv("data/iris.csv",na_values=['NA','?'])  
  
        encode_text_index(df,"species")    # encoding first before you call to_xy()  
  
        df
```

Out[8]:

	sepal_l	sepal_w	petal_l	petal_w	species
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0
5	5.4	3.9	1.7	0.4	0
6	4.6	3.4	1.4	0.3	0
7	5.0	3.4	1.5	0.2	0
8	4.4	2.9	1.4	0.2	0
9	4.9	3.1	1.5	0.1	0
10	5.4	3.7	1.5	0.2	0
11	4.8	3.4	1.6	0.2	0
12	4.8	3.0	1.4	0.1	0
13	4.3	3.0	1.1	0.1	0
14	5.8	4.0	1.2	0.2	0
15	5.7	4.4	1.5	0.4	0
16	5.4	3.9	1.3	0.4	0
17	5.1	3.5	1.4	0.3	0
18	5.7	3.8	1.7	0.3	0
19	5.1	3.8	1.5	0.3	0
20	5.4	3.4	1.7	0.2	0

## 16. Example of Deal with Missing values and outliers

```
In [11]: path = "./data/"  
  
        filename_read = os.path.join(path,"auto-mpg.csv")  
        df = pd.read_csv(filename_read,na_values=['NA','?'])  
  
        # Handle mising values in horsepower  
        missing_median(df, 'horsepower')  
        #df.drop('name', 1,inplace=True)  
  
        # Drop outliers in horsepower  
        print("Length before MPG outliers dropped: {}".format(len(df)))  
        remove_outliers(df,'mpg',2)  
        print("Length after MPG outliers dropped: {}".format(len(df)))  
  
        Length before MPG outliers dropped: 398  
        Length after MPG outliers dropped: 388
```

## 17. Training/Test split

Training data / Test Split means that the data are split according to a ratio with a training and validation set. Common ratios are 80% training and 20% validation.

```
In [30]: import pandas as pd
import io
import numpy as np
import os
from sklearn.model_selection import train_test_split

path = "./data/"

filename = os.path.join(path, "iris.csv")
df = pd.read_csv(filename, na_values=['NA', '?'])
df[0:5]
```

```
Out[30]:
```

	sepal_l	sepal_w	petal_l	petal_w	species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

```
In [31]: from sklearn import preprocessing

le = preprocessing.LabelEncoder()
df['encoded_species'] = le.fit_transform(df['species'])
df[0:5]
```

```
Out[31]:
```

	sepal_l	sepal_w	petal_l	petal_w	species	encoded_species
0	5.1	3.5	1.4	0.2	Iris-setosa	0
1	4.9	3.0	1.4	0.2	Iris-setosa	0
2	4.7	3.2	1.3	0.2	Iris-setosa	0
3	4.6	3.1	1.5	0.2	Iris-setosa	0
4	5.0	3.6	1.4	0.2	Iris-setosa	0

```
In [32]: # Split into train/test
x_train, x_test, y_train, y_test = train_test_split(df
```

```
In [33]: x_train.shape
```

```
Out[33]: (112, 4)
```

```
In [34]: y_train.shape
```

```
Out[34]: (112,)
```

```
In [35]: x_test.shape
```

```
Out[35]: (38, 4)
```


```
In [36]: y_test.shape
```

```
Out[36]: (38,)
```


This training data shows how the training set uses 80% of the data test validation set uses 20%.

## 18. Aggregation

- Data preprocessing where values of two or more objects are combined.
- Data aggregation can help reduce the size of data, change granularity of analysis, and improve stability of the data.
- Grouped by month



- Grouped by year



**19. Sampling:** 1, data reduction for exploratory data analysis and scaling up algorithms to big data applications and 2, quantify uncertainties due to varying data distributions.

- Sample size 3

```
In [42]: sample = data.sample(n=3)
sample

Out[42]:
```

	Clump Thickness	Uniformity of Cell Size	Uniformity of Cell Shape	Marginal Adhesion	Single Epithelial Cell Size	Bare Nuclei	Bland Chromatin	Normal Nucleoli	Mitoses	Class
77	5	3	1	2	2	1	2	1	1	2
511	5	1	1	1	2	1	2	1	1	2
519	4	7	8	3	4	10	9	1	1	4

- Randomly select 1% of data

```
In [43]: sample = data.sample(frac=0.01, random_state=1)
sample


Out[43]:
```

	Clump Thickness	Uniformity of Cell Size	Uniformity of Cell Shape	Marginal Adhesion	Single Epithelial Cell Size	Bare Nuclei	Bland Chromatin	Normal Nucleoli	Mitoses	Class
584	5	1	1	6	3	1	1	1	1	2
417	1	1	1	1	2	1	2	1	1	2
606	4	1	1	2	2	1	1	1	1	2
349	4	2	3	5	3	8	7	6	1	4
134	3	1	1	1	3	1	2	1	1	2
502	4	1	1	2	2	1	2	1	1	2
117	4	5	5	10	4	10	7	5	8	4

**20. Discretization:** data preprocessing step used to transform a continuous-valued attribute to a categorical attribute.

```
In [45]: data['Clump Thickness'].hist(bins=10)  
data['Clump Thickness'].value_counts(sort=False)
```

```
Out[45]: 1    145  
        2     50  
        3    108  
        4     80  
        5    130  
        6     34  
        7     23  
        8     46  
        9     14  
       10     69  
Name: Clump Thickness, dtype: int64
```



```
In [46]: bins = pd.cut(data['Clump Thickness'],4)  
bins.value_counts(sort=False)
```

```
Out[46]: (0.991, 3.25]    303  
        (3.25, 5.5]      210  
        (5.5, 7.75]       57  
        (7.75, 10.0]     129  
Name: Clump Thickness, dtype: int64
```

```
In [47]: bins = pd.qcut(data['Clump Thickness'],4)  
bins.value_counts(sort=False)
```

```
Out[47]: (0.999, 2.0]    195  
        (2.0, 4.0]      188  
        (4.0, 6.0]      164  
        (6.0, 10.0]     152  
Name: Clump Thickness, dtype: int64
```

**21. Principal Component Analysis:** PCA is a method for reducing the number of attributes in the data by projecting the data from its original high-dimensional space into a lower dimensional space. The new components have: 1. Linear combinations of the original attributes, 2. Perpendicular to each other, 3. They capture the maximum amount of variation in the data.

```
In [1]: %matplotlib inline
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
import numpy as np

numImages = 16
fig = plt.figure(figsize=(7,7))
imgData = np.zeros(shape=(numImages,36963))

for i in range(1,numImages+1):
    filename = 'pics/Picture'+str(i)+'.jpg'
    img = mpimg.imread(filename)
    ax = fig.add_subplot(4,4,i)
    plt.imshow(img)
    plt.axis('off')
    ax.set_title(str(i))
    imgData[i-1] = np.array(img.flatten()).reshape(1,img.shape[0]*img.shape[1]*img.shape[2])
```



## **Pre-Processing of our Data:**

From page 17 onward, we start the Pre-Processing of the datasets that we found online. Those datasets can be found on Page 2 of this document as hyperlinks.



## Missing Values - Talal (London Air)

```
Removing Missing Values

In [34]: import pandas as pd

data = pd.read_csv('C:\\Users\\Talal\\School\\CSC 177\\Labs\\Lab 1\\Missing Values\\LaqnData.csv')
data.columns = ['Site', 'Species', 'ReadingDate', 'Value', 'Units', 'Provisional or Ratified']

print('Number of instances = %d' % (data.shape[0]))
print('Number of attributes = %d' % (data.shape[1]))
data = data.drop(['Provisional or Ratified'], axis = 1)

#dropping all rows that contain missing values
data = data.dropna(axis='rows')

data.head()

#Writing modified data frame to csv file
data.to_csv(r'C:\\Users\\Talal\\School\\CSC 177\\Labs\\Lab 1\\Missing Values\\LaqnDataModified.csv')
#Resulting file has no missing value rows

Number of instances = 175200
Number of attributes = 6
```

### Original Dataset (with missing values):

In this dataset, multiple columns are left blank. The code above is meant to drop all of the columns that have missing values.

Site	Species	ReadingDate	Value	Units	Provisional or Ratified
H10	CO	#####		mg m-3	P
H10	CO	#####		mg m-3	P
H10	CO	#####		mg m-3	P
H10	CO	#####		mg m-3	P
H10	CO	#####		mg m-3	P

### Modified Dataset (with missing values and last column removed):

After running the code, the columns that were left blank were completely removed, leaving us with fewer columns to use for our analysis.

Site	Species	ReadingDate	Value	Units
H10	NO	#####	7.1	ug m-3
H10	NO	#####	7.1	ug m-3
H10	NO	#####	7.1	ug m-3
H10	NO	#####	7.1	ug m-3


## Outliers - Dianna

Outliers are data instances where the characteristics are notably different from the rest of the data. After running the code we can see that there were some outliers far above and below our data.

```
In [19]: %matplotlib inline
```

```
# data2 = data.drop(['Class'],axis=1)  
data.boxplot(figsize=(20,10))
```

```
Out[19]: <matplotlib.axes._subplots.AxesSubplot at 0x1212d94a8>
```



## Duplicate Data - Tom (Drugs)

- The `.drop_duplicates()` method will delete all duplicate rows, leaving us with a lot less rows to analyze in our data set.

```
Duplicate Data

In [4]: import os
import pandas as pd

path = "./ndcxls/"
filename_read = os.path.join(path, "product.csv")
df = pd.read_csv(filename_read, na_values=['NA', '?'])
df.drop(['PRODUCTID', 'PRODUCTNDC'], axis=1, inplace=True)
dups = df.duplicated()
print('Number of duplicate rows = %d' % (dups.sum()))
print('Number of rows before discarding duplicates = %d' % (df.shape[0]))
data2 = df.drop_duplicates()
print('Number of rows after discarding duplicates = %d' % (data2.shape[0]))

Number of duplicate rows = 5939
Number of rows before discarding duplicates = 129933
Number of rows after discarding duplicates = 123994
```

- After running the code, it clearly states that the number of rows before and after running the code went from “129933” to “123994”, which is about 6000 rows that we deleted. This is significant, because it is about 5% less data for us to look at after Pre-Processing.

## Shuffling Data Frames- Tom (Drugs)

- You can shuffle rows in your dataframe around in order to get a cleaner-looking dataset

```
# Shuffling Dataframes

In [27]: import os
import pandas as pd
import numpy as np

path = "./ndcxls/"
filename_read = os.path.join(path, "product.csv")
df = pd.read_csv(filename_read, na_values=['NA', '?'])
np.random.seed(38) # Uncomment this line to get the same shuffle each time

df = df.reindex(np.random.permutation(df.index))
df.reset_index(inplace=True, drop=True)
# use inplace=False
df

Out[27]:
```

	PRODUCTID	PRODUCTNDC	PRODUCTTYPE	PROPRIETARYNAME	PROPRIETARYNAME
0	0363-4200_4d13cab4-35ac-4407-8138-7bdd96344c84	0363-4200	HUMAN OTC DRUG	Daytime Nighttime Multi Symptom Cold	
1	63629-2598_3810a365-4f0f-4b39-9219-89f274534f6	63629-2598	HUMAN PRESCRIPTION DRUG	Nystatin	
2	50090-2914_2914_5f6b206d-8a4f-44a5-bd9d-85acd9c056f4	50090-2914	HUMAN PRESCRIPTION DRUG	quetiapine fumarate	
3	11716-1104_ddd93887-104b-4d41-8a4b-...	11716-1104	HUMAN OTC DRUG	LEADER ADVANCED FORMULA EYE DROPS	

- After shuffling, the output of the data is easier for the user to read and analyze

## Sorting Dataframes - Jimmy (Drugs)

- Columns can be sorted through the use of quicksort, mergesort, or heapsort. This is useful, because some data may be significant than other data for the purpose of analysis.

```
# Sorting Dataframes

In [33]: import os
import pandas as pd

path = "./ndcx1s/"
filename_read = os.path.join(path, "product.csv")
df = pd.read_csv(filename_read, na_values=['NA', '?'])

df = df.sort_values(by='PROPRIETARYNAME', ascending=True, axis=0)
df
```

Out[33]:

	PRODUCTID	PRODUCTNDC	PRODUCTTYPE	PROPRIETARYNAME	PROPRIETARYNAME
60653	51662-1260-7860dfb7-a67d-7b5f-e053-2a91aa0ad59f	51662-1260	HUMAN PRESCRIPTION DRUG	0.9% SODIUM CHLORIDE	
60671	51662-1301-78825011-268d-784b-e053-2991aa0ae0d4	51662-1301	HUMAN PRESCRIPTION DRUG	0.9% SODIUM CHLORIDE	
60675	51662-1308-7882aac1-44e8-cb42-e053-2991aa0a9a43	51662-1308	HUMAN PRESCRIPTION DRUG	0.9% SODIUM CHLORIDE	
74412	57691-161-369e0775-d4c4-404b-	57691-161	HUMAN OTC DRUG	01 Moisture	

- In this case, Proprietary Name was the column that was sorted, and this is because we were interested in the name of the chemical, rather than the actual ID, or Product Type

## Saving a Dataframe - Talal (London Air)

```
Saving a Dataframe

In [*]: import tkinter as tk
from tkinter import filedialog
import pandas as pd
global data
data = pd.read_csv('C:\\Users\\Talal\\School\\CSC 177\\Labs\\Lab 1\\Missing Values\\LaqnData.csv')
data.columns = ['Site', 'Species', 'ReadingDateTime', 'Value', 'Units', 'Provisional or Ratified']

print('Number of instances = %d' % (data.shape[0]))
print('Number of attributes = %d' % (data.shape[1]))
data = data.drop(['Provisional or Ratified'], axis = 1)

#dropping all rows that contain missing values
data = data.dropna(axis='rows')

data.head()

root = tk.Tk()

canvas1 = tk.Canvas(root, width = 300, height = 300, bg = 'lightsteelblue2', relief = 'raised')
canvas1.pack()

def saveCSV():
    exportPath = filedialog.asksavesavefilename(defaultextension = '.csv')
    data.to_csv(exportPath, index = None, header=True)

saveAsButton_CSV = tk.Button(text='Export CSV', command=saveCSV)
canvas1.create_window(150, 150, window=saveAsButton_CSV)

root.mainloop()

#This code needs to be run outside of Jupiter notebooks using Python through terminal
#to create a GUI window that will allow the user to save the modified (cleaned) dataset specified

Number of instances = 175200
Number of attributes = 6
```

## Dropping Fields - Jimmy (Drugs)

- You can drop multiple columns and rows with one line of code

### # Dropping Fields

```
In [44]: import os
import pandas as pd

path = "./ndcxls/"
filename_read = os.path.join(path, "product.csv")
df = pd.read_csv(filename_read, na_values=['NA', '?'])

df.drop(columns=['PRODUCTID', 'PRODUCTNDC'], index=[3], axis=1, inplace=True)
df[0:5]
```

Out[44]:

	PRODUCTTYPE	NAME	PROPRIETARYNAME	PROPRIETARYNAMESUFFIX	NONPROPRIETARYNAME	DOSAGE
0	HUMAN	OTC DRUG	Sterile Diluent	NaN	diluent	
1	HUMAN	PRESCRIPTION DRUG	Amyvid	NaN	Florbetapir F 18	
2	HUMAN	PRESCRIPTION DRUG	Quinidine Gluconate	NaN	Quinidine Gluconate	
4	HUMAN	PRESCRIPTION DRUG	Trulicity	NaN	Dulaglutide	
5	HUMAN	PRESCRIPTION DRUG	EMGALITY	NaN	galcanezumab	

- We chose to drop the “PRODUCTID”, and the “PRODUCTNDC” fields, because we deemed that they were unnecessary fields, and did not provide useful information.

## Calculated Fields - Amrit

- When doing calculations, all values must be numeric or you will get an error.

```
# Calculated Fields

In [74]: import os
import pandas as pd

path = "./ndcxls/"
filename_read = os.path.join(path, "product.csv")
df = pd.read_csv(filename_read, na_values=['NA', '?'])
#change none numeric values to NaN
df['ACTIVE_NUMERATOR_STRENGTH'] = pd.to_numeric(df['ACTIVE_NUMERATOR_STRENGTH'], errors='coerce')
#join 2 columns together
col_strength = df['ACTIVE_NUMERATOR_STRENGTH']
col_drugname = df['PROPRIETARYNAME']
result = pd.concat([col_drugname, col_strength], axis=1)
#drop all rows with NaN values
result = result.dropna()
#calculate
result.insert(1, '3x_strength', (result['ACTIVE_NUMERATOR_STRENGTH']*3).astype(float))
result
```

```
Out[74]:
```

	PROPRIETARYNAME	3x_strength	ACTIVE_NUMERATOR_STRENGTH
0	Sterile Diluent	3.00	1.00
1	Amyvid	153.00	51.00
2	Quinidine Gluconate	240.00	80.00
3	Trulicity	2.25	0.75
4	Trulicity	4.50	1.50
5	EMGALITY	360.00	120.00
6	TALTZ	240.00	80.00
7	EMGALITY	360.00	120.00
8	Prozac	270.00	90.00

- To ensure that all of the fields were numeric, we dropped the rows where there were “NaN” values.

## Feature Normalization

```
# Feature Normalization

In [79]: import os
import pandas as pd
import numpy as np
from scipy.stats import zscore

path = "./ndcxls/"
filename_read = os.path.join(path, "product.csv")
df = pd.read_csv(filename_read, na_values=['NA', '?'])
#change none numeric values to NaN
df['ACTIVE_NUMERATOR_STRENGTH'] = pd.to_numeric(df['ACTIVE_NUMERATOR_STRENGTH'], errors='coerce')
#join 2 columns together
col_strength = df['ACTIVE_NUMERATOR_STRENGTH']
col_drugname = df['PROPRIETARYNAME']
result = pd.concat([col_drugname, col_strength], axis=1)
#drop all rows with NaN values
result = result.dropna()
#calculate
result.insert(1, '3x_strength', (result['ACTIVE_NUMERATOR_STRENGTH']*3).astype(float))
#zscore
result['3x_strength'] = zscore(result['3x_strength'])
result
```

```
Out[79]:
```

	PROPRIETARYNAME	3x_strength	ACTIVE_NUMERATOR_STRENGTH
0	Sterile Diluent	-0.004984	1.00
1	Amyvid	-0.004982	51.00
2	Quinidine Gluconate	-0.004981	80.00
3	Trulicity	-0.004984	0.75
4	Trulicity	-0.004984	1.50
5	EMGALITY	-0.004980	120.00
6	TALTZ	-0.004981	80.00

## Concatenation - Min (Drugs)

- The method `.concat()` adds multiple columns of data together. Axis 1 adds columns horizontally and axis 0 adds everything in one column vertically.

```
In [2]: import os
import pandas as pd
import numpy as np
from scipy.stats import zscore

path = "./ndcx1s/"

filename_read = os.path.join(path, "product.csv")
df = pd.read_csv(filename_read, na_values=['NA', '?'])
col_strength = df['ACTIVE_NUMERATOR_STRENGTH']
col_drugname = df['PROPRIETARYNAME']
result = pd.concat([col_drugname, col_strength], axis=1)
result
```

Out[2]:

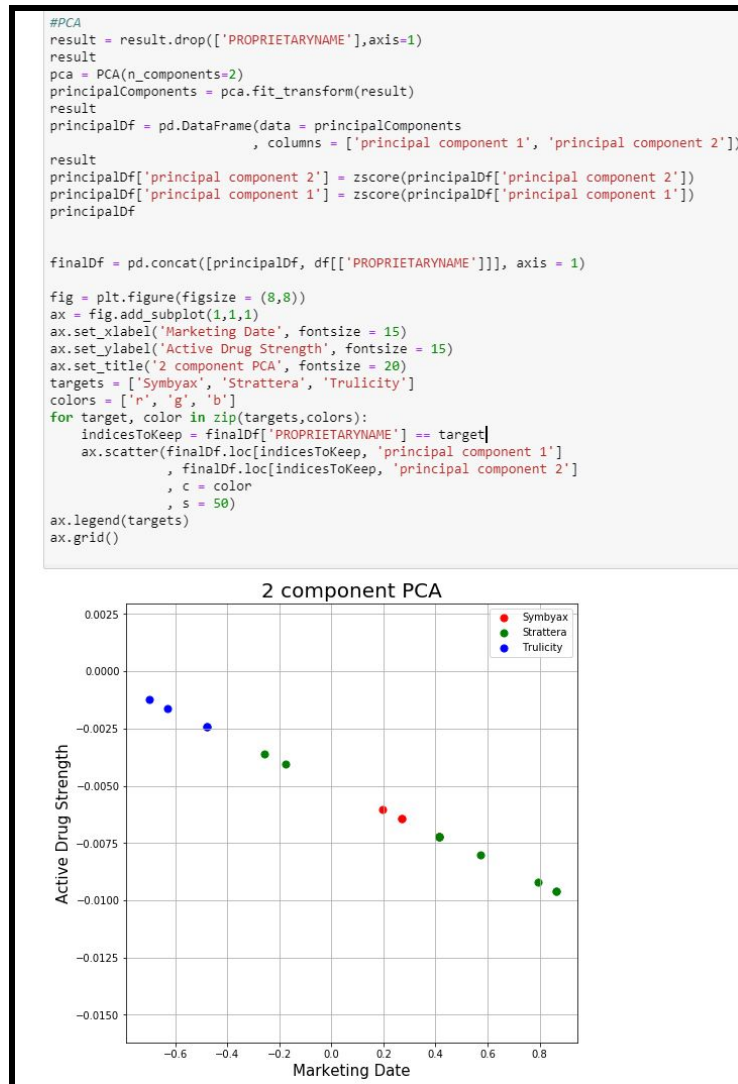
	PROPRIETARYNAME	ACTIVE_NUMERATOR_STRENGTH
0	Sterile Diluent	1
1	Amyvid	51
2	Quinidine Gluconate	80
3	Trulicity	0.75
4	Trulicity	1.5
5	EMGALITY	120
6	TALTZ	80
7	EMGALITY	120
8	Prozac	90
9	Strattera	10
10	Strattera	25
11	Strattera	40

## Aggregation

-

## PCA:

One thing that was interesting to analyze is the PCA performed on drugs Symbyax, Strattera, and Trulicity. PCA performed on active drug strength relative starting marketing date. From the PCA, the data suggests that as the Standardized (Z-Score) Marketing Date increases, the Strattera was manufactured at a much lower strength. When compared to Symbyax, Trulicity continued to stay unchanged. It was interesting to see a negative correlation in a large set of data.





## Data split

```
#---Training/Test split---
filename_read = os.path.join(path, 'product2.csv')
df = pd.read_csv(filename_read, na_values=['NA', '?'])

le = preprocessing.LabelEncoder()
df['encoded_species'] = le.fit_transform(df['Species'])

x_train, x_test, y_train, y_test = train_test_split(df[['Species', 'Value']],
                                                    df['encoded_species'],
                                                    test_size=0.25, random_state=60)

print('Training set mean = %d' % (x_train.Value.mean()))
print('Training set standard deviation = %d' % (x_train.Value.std()))
print('Test set mean = %d' % (x_test.Value.mean()))
print('Test set standard deviation = %d' % (x_test.Value.std()))
```


```
Training set mean = 47
Training set standard deviation = 54
Test set mean = 48
Test set standard deviation = 54
```

---

## Aggregation

```
#---AGGREGATION---
daily = df_air
daily.index = pd.to_datetime(daily['ReadingDateTime'])
daily = daily['Value']
fig = plt.figure(figsize = (8,8))
ax1 = fig.add_subplot(1,1,1)
ax1 = daily.plot(kind = 'line',figsize=(15,3))
ax1.set_title('Daily Air Quality (variance = %.4f)' % (daily.var()))
```

Out[40]: Text(0.5, 1.0, 'Daily Air Quality (variance = 2943.2929)')



### 1) Interpreting human insights on the data and its possible effect on predictions

Through the Pre-Processing of our data, the number of rows and columns were significantly reduced. Since we “scrubbed” the data so that it became cleaner, our confidence in the legitimacy of the data increased. We got rid of NaN data, blank data, unnecessary fields, duplicate data, outliers, and sorted the data in such a way where the format was easily readable. After condensing the data into what we deemed necessary, the amount of information to look at is significantly less for us to read, and less for our program to parse through. This will result in quicker, and more accurate predictions, because of the removal of junk data and dimensions.

### 2) Split data into two sets: training data (first 80% of the product.csv file), and test data (last 20% of the product.csv file), and calculate the mean and standard deviation

#### **Training Data:**

Mean: 47

Standard Deviation: 54

#### **Test Data:**

Mean: 48

Standard Deviation: 54

### 3) Document ratios for splitting and the results

We decided to primarily focus on the “Drugs” link, specifically “Product.csv”. The reason for this is because we have done the most thorough Pre-Processing on this data set. We decided to split the data as follows:

Training Data: First 80% of the data after Pre-Processing

Test Data: Last 20% of the data after Pre-Processing

The reason we decided on 80/20 is because since we did “Sorting Dataframe”, the rows are all randomized. That’s why it sufficient to pick the first 80%, and the last 20% of the data.

Results:

- 4) Compare the two sets: the training data and the test data and analyze it, developing an intuition and meaning of your results.

What we can draw from this data is that splitting data into two consequent sets is a vital step in creating a machine learning model. The training data receives the majority share of the partitions since this is the data that the model is built on, whereas the remaining test set will be used to validate the training data against. In our case, it is apparent that the two sets of randomized data provide consistent results in terms of the mean and standard deviation of the “value” attribute.